

```
In [2]: import pandas as pd
import numpy as np
import re
import string
import nltk
from nltk.corpus import stopwords
from nltk.tokenize import word_tokenize
from nltk.stem import WordNetLemmatizer
import matplotlib.pyplot as plt
import seaborn as sns
from collections import Counter
import warnings

warnings.filterwarnings('ignore')
plt.style.use('seaborn-v0_8-darkgrid')

print("Downloading required NLTK data...")
nltk.download('punkt', quiet=True)
nltk.download('stopwords', quiet=True)
nltk.download('wordnet', quiet=True)
nltk.download('averaged_perceptron_tagger', quiet=True)
nltk.download('omw-1.4', quiet=True)

print("✅ All libraries imported and NLTK data downloaded successfully!")
```

Downloading required NLTK data...

✅ All libraries imported and NLTK data downloaded successfully!

```
In [3]: data_path = '../data/processed/merged_health_data.csv'

df = pd.read_csv(data_path)

print("*"*60)
print("DATA LOADED SUCCESSFULLY")
print("*"*60)
print(f"Total Records: {len(df)}")
print(f"Columns: {list(df.columns)}")
print(f"\nFirst 3 rows:")
print(df[['title', 'text', 'health_category']].head(3))
```

```
=====
DATA LOADED SUCCESSFULLY
=====
Total Records: 19,080
Columns: ['id', 'title', 'text', 'created_utc', 'score', 'num_comments', 'subredditt', 'category', 'health_category', 'source_file']

First 3 rows:
      title                                     text \
0  Diabetes and wounds  35. M. I've been a bad diabetic for a while, a...
1      T1D in Movies  Curious, for the movie buffs on the sub, what ...
2          Patches  Hi, I'm caring for someone who has very fragil...
```

	health_category	
0	diabetes	
1	diabetes	
2	diabetes	

```
In [4]: lemmatizer = WordNetLemmatizer()

stop_words = set(stopwords.words('english'))

custom_stopwords = {
    'im', 'ive', 'dont', 'didnt', 'doesnt', 'isnt', 'arent', 'wasnt', 'werent',
    'hasnt', 'havent', 'hadnt', 'wont', 'wouldnt', 'couldnt', 'shouldnt',
    'cant', 'cannot', 'mightnt', 'mustnt', 'shant', 'shan',
    'id', 'youre', 'youve', 'youll', 'youd', 'hes', 'shes', 'its',
    'were', 'theyve', 'theyll', 'theyd', 'whats', 'wheres', 'whens',
    'whos', 'whys', 'hows', 'theres', 'heres',
    'reddit', 'post', 'edit', 'update', 'tldr', 'amp', 'x200b'
}

stop_words.update(custom_stopwords)

print("✅ Preprocessing tools initialized!")
print(f"Total stopwords: {len(stop_words)}")
```

✅ Preprocessing tools initialized!

Total stopwords: 244

```
In [5]: def clean_text(text):
    if pd.isna(text) or text == '':
        return ''

    text = str(text).lower()

    text = re.sub(r'\[deleted\]|\[removed\]', '', text)

    text = re.sub(r'http\S+|www\S+|https\S+', '', text, flags=re.MULTILINE)

    text = re.sub(r'@\w+', '', text)

    text = re.sub(r'#\w+', '', text)
```

```

text = re.sub(r'\n', ' ', text)
text = re.sub(r'\t', ' ', text)
text = re.sub(r'\r', ' ', text)

text = re.sub(r'&|<|>|"', '', text)

text = re.sub(r'\d+', '', text)

text = text.translate(str.maketrans(' ', ' ', string.punctuation))

text = re.sub(r'\s+', ' ', text).strip()

return text

def tokenize_text(text):
    if text == '':
        return []

    try:
        tokens = word_tokenize(text)
    except:
        tokens = text.split()

    return tokens

def remove_stopwords(tokens):
    filtered_tokens = [word for word in tokens if word not in stop_words and len(word) > 1]
    return filtered_tokens

def lemmatize_tokens(tokens):
    lemmatized = [lemmatizer.lemmatize(word) for word in tokens]
    return lemmatized

def preprocess_text(text):
    cleaned = clean_text(text)

    tokens = tokenize_text(cleaned)

    tokens = remove_stopwords(tokens)

    tokens = lemmatize_tokens(tokens)

    processed_text = ' '.join(tokens)

    return processed_text

print("✅ Text cleaning functions defined!")
print("\nTesting on sample text:")
sample = "I'm really struggling with my diabetes! Check out this link: https://www.diabetessucks.com"
print(f"Original: {sample}")
print(f"Cleaned: {preprocess_text(sample)}")

```

✓ Text cleaning functions defined!

Testing on sample text:

Original: I'm really struggling with my diabetes! Check out this link: http

s://example.com #health @doctor123

Cleaned: really struggling diabetes check link

```
In [6]: print("*"*60)
print("STARTING TEXT PREPROCESSING")
print("*"*60)
print("This may take several minutes for large datasets...")
print()

df['cleaned_text'] = df['text'].apply(preprocess_text)

df['cleaned_title'] = df['title'].apply(preprocess_text)

df['combined_text'] = df['cleaned_title'] + ' ' + df['cleaned_text']

df['combined_text'] = df['combined_text'].str.strip()

print("✓ Preprocessing completed!")
print(f"\nSample of cleaned data:")
print("-" * 60)

for idx in df.sample(3).index:
    print(f"\nOriginal: {df.loc[idx, 'text'][:150]}...")
    print(f"Cleaned: {df.loc[idx, 'cleaned_text'][:150]}...")
    print("-" * 60)
```

```
=====
```

STARTING TEXT PREPROCESSING

```
=====
```

This may take several minutes for large datasets...

✓ Preprocessing completed!

Sample of cleaned data:

```
-----
```

Original: raw shredded beets, shredded carrots, grilled chicken, cottage cheese, red onions, sautéed bell peppers & broccoli, baked sweet potato. drizzle of pomegranate molasses...

Cleaned: raw shredded beet shredded carrot grilled chicken cottage cheese red onion sautéed bell pepper broccoli baked sweet potato drizzle pomegranate molasses...

```
-----
```

Original: What are my options for addressing my narrow airway while I am awake? I have issues breathing and doing things like cardio because of my narrow airway...

Cleaned: option addressing narrow airway awake issue breathing thing like cardio narrow airway nice addressed sleeping like address waking hour...

```
-----
```

Original: My personal go to tip if experiencing panic attacks or feel it starting is to go to the bathroom sink and run cold water over hands and face or get in...

Cleaned: personal tip experiencing panic attack feel starting bathroom sink run cold water hand face get freezing cold shower minute help shock brain back real...

```
-----
```

```
In [7]: print("=*60)
print("POST-PREPROCESSING STATISTICS")
print("=*60)

df['cleaned_length'] = df['cleaned_text'].apply(len)
df['cleaned_word_count'] = df['cleaned_text'].apply(lambda x: len(x.split()))

df['original_length'] = df['text'].fillna('').astype(str).apply(len)
df['original_word_count'] = df['text'].fillna('').astype(str).apply(lambda x:

print(f"\nOriginal Data:")
print(f"  Average text length: {df['original_length'].mean():.0f} characters")
print(f"  Average word count: {df['original_word_count'].mean():.0f} words")

print(f"\nCleaned Data:")
print(f"  Average text length: {df['cleaned_length'].mean():.0f} characters")
print(f"  Average word count: {df['cleaned_word_count'].mean():.0f} words")

reduction_chars = ((df['original_length'].mean() - df['cleaned_length'].mean())
reduction_words = ((df['original_word_count'].mean() - df['cleaned_word_count'].mean())
```

```

print(f"\nReduction:")
print(f"  Characters: {reduction_chars:.1f}%")
print(f"  Words: {reduction_words:.1f}%")

empty_after_cleaning = len(df[df['cleaned_text'] == ''])
print(f"\nEmpty texts after cleaning: {empty_after_cleaning} ({empty_after_cleaning * 100 / len(df)}%")

```

=====

POST-PREPROCESSING STATISTICS

=====

Original Data:

Average text length: 812 characters
 Average word count: 145 words

Cleaned Data:

Average text length: 477 characters
 Average word count: 71 words

Reduction:

Characters: 41.3%
 Words: 51.2%

Empty texts after cleaning: 16 (0.08%)

```
In [8]: fig, axes = plt.subplots(2, 2, figsize=(16, 10))

axes[0, 0].hist(df['original_word_count'], bins=50, alpha=0.7, color='red', label='Original')
axes[0, 0].hist(df['cleaned_word_count'], bins=50, alpha=0.7, color='green', label='Cleaned')
axes[0, 0].set_xlabel('Word Count', fontsize=12)
axes[0, 0].set_ylabel('Frequency', fontsize=12)
axes[0, 0].set_title('Word Count Distribution: Before vs After', fontsize=14, fontweight='bold')
axes[0, 0].legend()
axes[0, 0].set_xlim(0, 500)

word_count_by_category = df.groupby('health_category')['cleaned_word_count'].mean()
axes[0, 1].barh(range(len(word_count_by_category)), word_count_by_category.values,
                 color=sns.color_palette("husl", len(word_count_by_category)))
axes[0, 1].set_yticks(range(len(word_count_by_category)))
axes[0, 1].set_yticklabels(word_count_by_category.index)
axes[0, 1].set_xlabel('Average Word Count', fontsize=12)
axes[0, 1].set_title('Average Cleaned Word Count by Category', fontsize=14, fontweight='bold')
axes[0, 1].invert_yaxis()

df_non_empty = df[df['cleaned_text'] != '']
axes[1, 0].boxplot([df_non_empty[df_non_empty['health_category'] == cat]['cleaned_text']
                    for cat in df['health_category'].unique()],
                   labels=df['health_category'].unique())
axes[1, 0].set_xlabel('Health Category', fontsize=12)
axes[1, 0].set_ylabel('Word Count', fontsize=12)
axes[1, 0].set_title('Word Count Distribution by Category', fontsize=14, fontweight='bold')
axes[1, 0].xaxis.set_tick_params(rotation=45)

reduction_per_category = df.groupby('health_category').apply(lambda x: {
    'category': x['health_category'].values[0],
    'reduction_chars': reduction_chars[x['health_category'].values[0]],
    'reduction_words': reduction_words[x['health_category'].values[0]]})
```

```

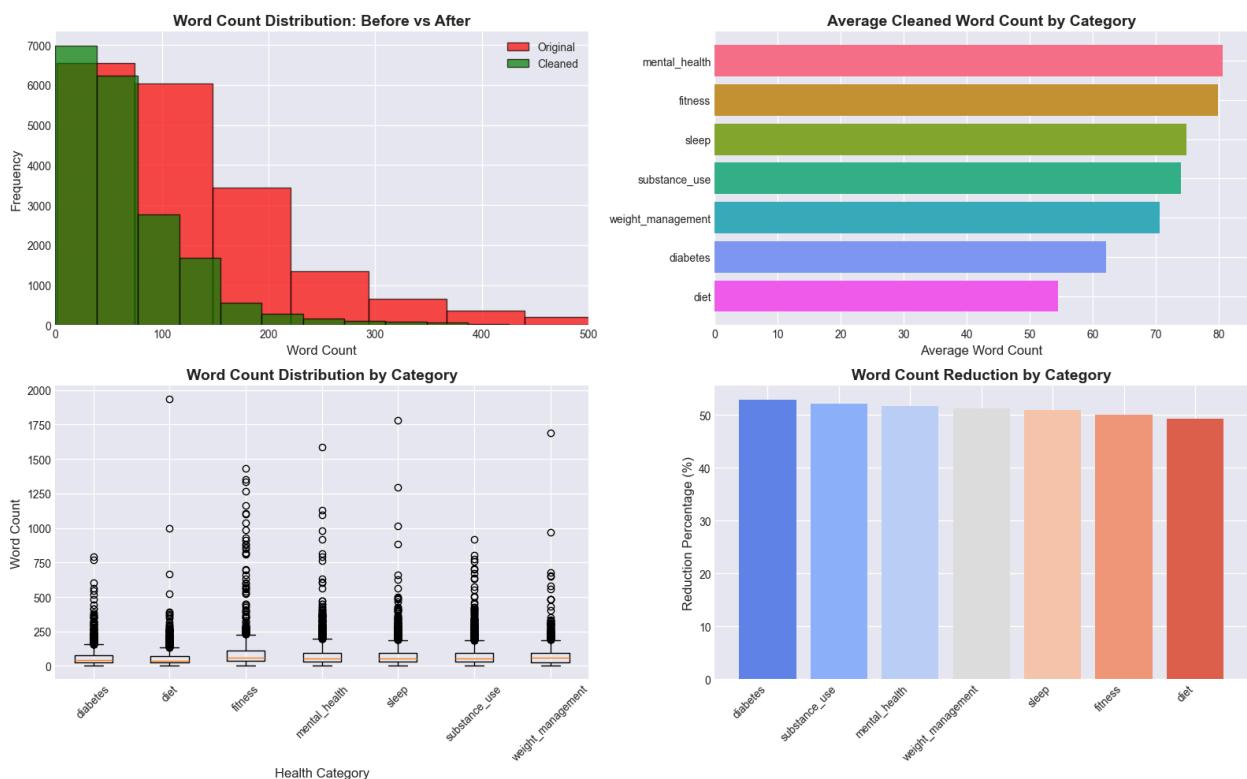
    lambda x: ((x['original_word_count'].mean() - x['cleaned_word_count'].mean())
    .sort_values(ascending=False))

    axes[1, 1].bar(range(len(reduction_per_category)), reduction_per_category.values,
                  color=sns.color_palette("coolwarm", len(reduction_per_category)))
    axes[1, 1].set_xticks(range(len(reduction_per_category)))
    axes[1, 1].set_xticklabels(reduction_per_category.index, rotation=45, ha='right')
    axes[1, 1].set_ylabel('Reduction Percentage (%)', fontsize=12)
    axes[1, 1].set_title('Word Count Reduction by Category', fontsize=14, fontweight='bold')

    plt.tight_layout()
    plt.savefig('../reports/preprocessing_impact.png', dpi=300, bbox_inches='tight')
    plt.show()

print("✅ Visualization saved to: reports/preprocessing_impact.png")

```



✅ Visualization saved to: reports/preprocessing_impact.png

```

In [9]: print("*"*60)
print("VOCABULARY ANALYSIS")
print("*"*60)

all_words = []
for text in df['cleaned_text']:
    if text != '':
        all_words.extend(text.split())

word_freq = Counter(all_words)
total_words = len(all_words)
unique_words = len(word_freq)

```

```

print(f"\nTotal words (after cleaning): {total_words}")
print(f"Unique words (vocabulary size): {unique_words}")
print(f"Average word frequency: {total_words/unique_words:.2f}")

print(f"\nTop 30 most common words:")
print("-" * 60)
for word, count in word_freq.most_common(30):
    percentage = (count / total_words) * 100
    print(f"{word:20s}: {count}, ({percentage:5.2f}%)")

```

=====

VOCABULARY ANALYSIS

=====

Total words (after cleaning): 1,354,316
 Unique words (vocabulary size): 40,800
 Average word frequency: 33.19

Top 30 most common words:

like	:	14,111	(1.04%)
day	:	12,821	(0.95%)
i'm	:	10,992	(0.81%)
time	:	9,947	(0.73%)
get	:	9,568	(0.71%)
feel	:	8,858	(0.65%)
sleep	:	8,198	(0.61%)
year	:	7,626	(0.56%)
know	:	7,278	(0.54%)
also	:	6,952	(0.51%)
question	:	6,553	(0.48%)
week	:	6,320	(0.47%)
would	:	6,222	(0.46%)
one	:	6,199	(0.46%)
really	:	5,856	(0.43%)
i've	:	5,567	(0.41%)
thing	:	5,528	(0.41%)
want	:	5,473	(0.40%)
even	:	5,225	(0.39%)
make	:	5,160	(0.38%)
back	:	5,087	(0.38%)
month	:	4,864	(0.36%)
it's	:	4,848	(0.36%)
anyone	:	4,791	(0.35%)
still	:	4,786	(0.35%)
help	:	4,672	(0.34%)
much	:	4,654	(0.34%)
work	:	4,567	(0.34%)
weight	:	4,450	(0.33%)
good	:	4,244	(0.31%)

In [10]: `print("=*60)`
`print("VOCABULARY ANALYSIS BY CATEGORY")`
`print("=*60)`

```
category_vocab = {}

for category in df['health_category'].unique():
    category_text = ' '.join(df[df['health_category'] == category]['cleaned_text'])
    words = category_text.split()
    category_vocab[category] = {
        'total_words': len(words),
        'unique_words': len(set(words)),
        'top_words': Counter(words).most_common(10)
    }

for category, stats in category_vocab.items():
    print(f"\n{category.upper()}:")
    print(f"  Total words: {stats['total_words']},")
    print(f"  Unique words: {stats['unique_words']},")
    print(f"  Vocabulary richness: {stats['unique_words']/stats['total_words']}")
    print(f"  Top 10 words:")
    for word, count in stats['top_words']:
        print(f"    {word:15s}: {count:5,}")
```

=====

VOCABULARY ANALYSIS BY CATEGORY

=====

DIABETES:

Total words: 157,389
Unique words: 12,104
Vocabulary richness: 7.69%

Top 10 words:

like	:	1,388
sugar	:	1,322
i'm	:	1,310
day	:	1,253
blood	:	1,077
insulin	:	1,060
get	:	1,059
year	:	999
time	:	998
know	:	969

DIET:

Total words: 156,558
Unique words: 12,297
Vocabulary richness: 7.85%

Top 10 words:

like	:	1,605
keto	:	1,598
day	:	1,315
food	:	1,163
eat	:	1,153
i'm	:	1,102
protein	:	966
time	:	861
meal	:	849
get	:	847

FITNESS:

Total words: 231,052
Unique words: 10,448
Vocabulary richness: 4.52%

Top 10 words:

question	:	5,422
thread	:	2,628
like	:	2,267
also	:	2,243
search	:	2,074
rfitness	:	2,000
sure	:	1,927
get	:	1,637
training	:	1,483
fitness	:	1,408

MENTAL_HEALTH:

Total words: 218,914

Unique words: 14,278
Vocabulary richness: 6.52%
Top 10 words:

pain	:	2,938
like	:	2,403
i'm	:	1,922
migraine	:	1,894
get	:	1,731
day	:	1,645
feel	:	1,629
anxiety	:	1,535
time	:	1,429
know	:	1,285

SLEEP:

Total words: 207,248
Unique words: 13,591
Vocabulary richness: 6.56%
Top 10 words:

sleep	:	6,914
like	:	2,492
night	:	2,452
i'm	:	1,852
get	:	1,780
feel	:	1,764
time	:	1,727
hour	:	1,694
day	:	1,665
asleep	:	1,329

SUBSTANCE_USE:

Total words: 204,854
Unique words: 13,418
Vocabulary richness: 6.55%
Top 10 words:

day	:	3,412
i'm	:	2,302
like	:	2,250
time	:	2,246
feel	:	2,222
year	:	2,120
get	:	1,442
smoking	:	1,415
quit	:	1,368
know	:	1,309

WEIGHT_MANAGEMENT:

Total words: 178,301
Unique words: 11,557
Vocabulary richness: 6.48%
Top 10 words:

weight	:	2,411
day	:	2,202
i'm	:	1,788

```
like          : 1,706
calorie       : 1,662
time          : 1,303
eat           : 1,175
week          : 1,153
eating         : 1,144
feel           : 1,088
```

```
In [11]: print("*60")
print("FINAL DATA CLEANING")
print("*60)

original_count = len(df)

df_clean = df[df['cleaned_text'] != ''].copy()

df_clean = df_clean[df_clean['cleaned_word_count'] >= 3].copy()

removed_count = original_count - len(df_clean)

print(f"Original dataset: {original_count:,} records")
print(f"After removing empty/short texts: {len(df_clean):,} records")
print(f"Removed: {removed_count:,} records ({removed_count/original_count*100:.2f}%)")

print("\nFinal dataset distribution:")
print(df_clean['health_category'].value_counts())
```

```
=====
FINAL DATA CLEANING
=====
```

```
Original dataset: 19,080 records
After removing empty/short texts: 19,060 records
Removed: 20 records (0.10%)
```

```
Final dataset distribution:
health_category
fitness          2892
diet             2874
sleep            2766
substance_use   2763
mental_health    2715
diabetes         2531
weight_management 2519
Name: count, dtype: int64
```

```
In [12]: print("*60")
print("TEXT QUALITY METRICS")
print("*60)

def calculate_lexical_diversity(text):
    words = text.split()
    if len(words) == 0:
        return 0
    return len(set(words)) / len(words)
```

```

df_clean['lexical_diversity'] = df_clean['cleaned_text'].apply(calculate_lexic)

print(f"\nLexical Diversity Statistics:")
print(f"  Mean: {df_clean['lexical_diversity'].mean():.3f}")
print(f"  Median: {df_clean['lexical_diversity'].median():.3f}")
print(f"  Std Dev: {df_clean['lexical_diversity'].std():.3f}")

print(f"\nBy Category:")
diversity_by_category = df_clean.groupby('health_category')['lexical_diversity']
for category, diversity in diversity_by_category.items():
    print(f"  {category:25s}: {diversity:.3f}")

plt.figure(figsize=(12, 6))
plt.bar(range(len(diversity_by_category)), diversity_by_category.values,
        color=sns.color_palette("viridis", len(diversity_by_category)))
plt.xticks(range(len(diversity_by_category)), diversity_by_category.index, rot=90)
plt.ylabel('Lexical Diversity', fontsize=12)
plt.xlabel('Health Category', fontsize=12)
plt.title('Lexical Diversity by Health Category', fontsize=14, fontweight='bold')
plt.axhline(y=df_clean['lexical_diversity'].mean(), color='red', linestyle='--',
            label=f'Overall Mean: {df_clean["lexical_diversity"].mean():.3f}')
plt.legend()
plt.tight_layout()
plt.savefig('../reports/lexical_diversity.png', dpi=300, bbox_inches='tight')
plt.show()

print("\n✓ Visualization saved to: reports/lexical_diversity.png")

```

=====

TEXT QUALITY METRICS

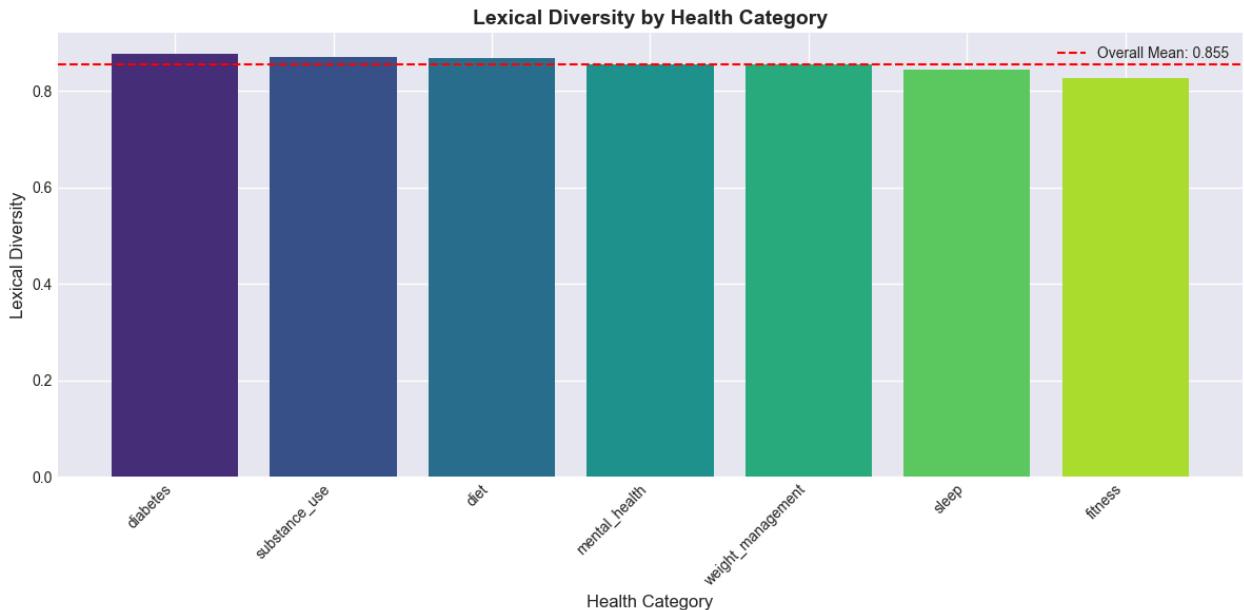
=====

Lexical Diversity Statistics:

Mean: 0.855
 Median: 0.857
 Std Dev: 0.094

By Category:

diabetes	:	0.876
substance_use	:	0.869
diet	:	0.869
mental_health	:	0.855
weight_management	:	0.854
sleep	:	0.843
fitness	:	0.825



✓ Visualization saved to: reports/lexical_diversity.png

```
In [13]: print("*"*60)
print("CREATING ADDITIONAL FEATURES")
print("*"*60)

df_clean['text_cleaned_length'] = df_clean['cleaned_text'].apply(len)

df_clean['avg_word_length'] = df_clean['cleaned_text'].apply(
    lambda x: np.mean([len(word) for word in x.split()]) if x else 0
)

df_clean['sentence_count'] = df_clean['text'].fillna('').apply(
    lambda x: len([s for s in re.split(r'[.!?]+', str(x)) if s.strip()])
)

df_clean['has_question'] = df_clean['text'].fillna('').str.contains(r'\?', re.IGNORECASE)

df_clean['has_exclamation'] = df_clean['text'].fillna('').str.contains(r'!', re.IGNORECASE)

df_clean['uppercase_ratio'] = df_clean['text'].fillna('').apply(
    lambda x: sum(1 for c in str(x) if c.isupper()) / len(str(x)) if len(str(x)) > 0 else 0
)

print("✓ Additional features created!")
print("\nNew feature columns:")
feature_cols = ['text_cleaned_length', 'avg_word_length', 'sentence_count',
                'has_question', 'has_exclamation', 'uppercase_ratio', 'lexical_diversity']
print(feature_cols)

print("\nSample statistics:")
print(df_clean[feature_cols].describe())
```

```
=====
CREATING ADDITIONAL FEATURES
=====
```

✓ Additional features created!

New feature columns:

```
['text_cleaned_length', 'avg_word_length', 'sentence_count', 'has_question', 'has_exclamation', 'uppercase_ratio', 'lexical_diversity']
```

Sample statistics:

	text_cleaned_length	avg_word_length	sentence_count	has_question	\
count	19060.000000	19060.000000	19060.000000	19060.000000	
mean	477.367681	5.696428	10.242078	0.602886	
std	545.198123	0.504391	11.469455	0.489313	
min	16.000000	3.750000	1.000000	0.000000	
25%	187.000000	5.362603	4.000000	0.000000	
50%	337.000000	5.655172	7.000000	1.000000	
75%	599.000000	6.000000	12.000000	1.000000	
max	13639.000000	9.666667	224.000000	1.000000	

	has_exclamation	uppercase_ratio	lexical_diversity
count	19060.000000	19060.000000	19060.000000
mean	0.261700	0.025792	0.855407
std	0.439572	0.017342	0.094427
min	0.000000	0.000000	0.136761
25%	0.000000	0.017937	0.790698
50%	0.000000	0.023443	0.857143
75%	1.000000	0.030534	0.923077
max	1.000000	0.788732	1.000000

```
In [14]: print("*60)
print("SAVING PREPROCESSED DATA")
print("*60)

columns_to_save = [
    'id', 'title', 'text', 'cleaned_title', 'cleaned_text', 'combined_text',
    'health_category', 'subreddit', 'score', 'num_comments', 'created_utc',
    'cleaned_word_count', 'text_cleaned_length', 'avg_word_length',
    'sentence_count', 'has_question', 'has_exclamation',
    'uppercase_ratio', 'lexical_diversity'
]

available_columns = [col for col in columns_to_save if col in df_clean.columns]

df_output = df_clean[available_columns].copy()

output_path = '../data/processed/preprocessed_data.csv'
df_output.to_csv(output_path, index=False)

print(f"✓ Preprocessed data saved to: {output_path}")
print(f"\nDataset shape: {df_output.shape}")
print(f"Columns saved: {len(available_columns)}")
print(f"\nColumn list:")
for col in available_columns:
```

```
    print(f" - {col}")

=====
SAVING PREPROCESSED DATA
=====
✓ Preprocessed data saved to: ./data/processed/preprocessed_data.csv

Dataset shape: (19060, 19)
Columns saved: 19

Column list:
- id
- title
- text
- cleaned_title
- cleaned_text
- combined_text
- health_category
- subreddit
- score
- num_comments
- created_utc
- cleaned_word_count
- text_cleaned_length
- avg_word_length
- sentence_count
- has_question
- has_exclamation
- uppercase_ratio
- lexical_diversity
```

```
In [15]: print("*60)
print("SAVING CATEGORY-WISE DATA FILES")
print("*60)

category_output_dir = '../data/processed/by_category/'

import os
os.makedirs(category_output_dir, exist_ok=True)

for category in df_clean['health_category'].unique():
    category_df = df_clean[df_clean['health_category'] == category][available_]

    category_file = f'{category_output_dir}{category}_preprocessed.csv'
    category_df.to_csv(category_file, index=False)

    print(f"✓ Saved {category}: {len(category_df)} records → {category_file}")

print("\n✓ All category files saved to: {category_output_dir}")
```

```
=====
SAVING CATEGORY-WISE DATA FILES
=====

✓ Saved diabetes: 2,531 records → ../data/processed/by_category/diabetes_preprocessed.csv
✓ Saved diet: 2,874 records → ../data/processed/by_category/diet_preprocessed.csv
✓ Saved fitness: 2,892 records → ../data/processed/by_category/fitness_preprocessed.csv
✓ Saved mental_health: 2,715 records → ../data/processed/by_category/mental_health_preprocessed.csv
✓ Saved sleep: 2,766 records → ../data/processed/by_category/sleep_preprocessed.csv
✓ Saved substance_use: 2,763 records → ../data/processed/by_category/substance_use_preprocessed.csv
✓ Saved weight_management: 2,519 records → ../data/processed/by_category/weight_management_preprocessed.csv

✓ All category files saved to: ../data/processed/by_category/
```

```
In [16]: print("=*80")
print(" " * 25 + "PREPROCESSING SUMMARY REPORT")
print("=*80")

print("\n📊 DATA TRANSFORMATION:")
print("-" * 80)
print(f" Original dataset: {original_count:,} records")
print(f" Final dataset: {len(df_clean):,} records")
print(f" Records removed: {removed_count:,} ({removed_count/original_count*100:.2f}%)")

print("\n📝 TEXT STATISTICS:")
print("-" * 80)
print(f" Average cleaned text length: {df_clean['text_cleaned_length'].mean()}")
print(f" Average cleaned word count: {df_clean['cleaned_word_count'].mean():.2f}")
print(f" Average word length: {df_clean['avg_word_length'].mean():.2f} characters")
print(f" Average lexical diversity: {df_clean['lexical_diversity'].mean():.3f}")

print("\n📚 VOCABULARY:")
print("-" * 80)
print(f" Total words (corpus): {total_words:,}")
print(f" Unique words (vocabulary): {unique_words:,}")
print(f" Vocabulary richness: {unique_words/total_words*100:.2f}%")

print("\n📈 CATEGORY DISTRIBUTION (FINAL):")
print("-" * 80)
for category, count in df_clean['health_category'].value_counts().items():
    percentage = (count / len(df_clean)) * 100
    print(f" {category:25s}: {count:6,} records ({percentage:5.2f}%)")

print("\n💡 FEATURE ENGINEERING:")
print("-" * 80)
print(f" Features created: {len(feature_cols)}")
print(" Features list:")
for feat in feature_cols:
```

```
print(f"      - {feat}")

print("\n📁 OUTPUT FILES:")
print("-" * 80)
print(f"  ✓ Main file: {output_path}")
print(f"  ✓ Category files: {len(df_clean['health_category'].unique())} files")
print(f"  ✓ Visualizations:")
print(f"    - preprocessing_impact.png")
print(f"    - lexical_diversity.png")

print("\n⌚ PREPROCESSING PIPELINE STEPS COMPLETED:")
print("-" * 80)
print("  ✓ Text cleaning (URLs, special chars, punctuation)")
print("  ✓ Tokenization")
print("  ✓ Stopword removal (English + custom)")
print("  ✓ Lemmatization")
print("  ✓ Empty text removal")
print("  ✓ Feature engineering")
print("  ✓ Data validation")
print("  ✓ Category-wise file generation")

print("\n" + "="*80)
print(" " * 25 + "PREPROCESSING COMPLETED SUCCESSFULLY! ✓")
print("=*80")
```

=====
=
PREPROCESSING SUMMARY REPORT
=====
=

 DATA TRANSFORMATION:

-
Original dataset: 19,080 records
Final dataset: 19,060 records
Records removed: 20 (0.10%)

 TEXT STATISTICS:

-
Average cleaned text length: 477 characters
Average cleaned word count: 71 words
Average word length: 5.70 characters
Average lexical diversity: 0.855

 VOCABULARY:

-
Total words (corpus): 1,354,316
Unique words (vocabulary): 40,800
Vocabulary richness: 3.01%

 CATEGORY DISTRIBUTION (FINAL):

-
fitness : 2,892 records (15.17%)
diet : 2,874 records (15.08%)
sleep : 2,766 records (14.51%)
substance_use : 2,763 records (14.50%)
mental_health : 2,715 records (14.24%)
diabetes : 2,531 records (13.28%)
weight_management : 2,519 records (13.22%)

 FEATURE ENGINEERING:

-
Features created: 7
Features list:
- text_cleaned_length
- avg_word_length
- sentence_count
- has_question
- has_exclamation
- uppercase_ratio
- lexical_diversity

 OUTPUT FILES:

- ✓ Main file: ../data/processed/preprocessed_data.csv
- ✓ Category files: 7 files in ../data/processed/by_category/
- ✓ Visualizations:
 - preprocessing_impact.png
 - lexical_diversity.png

⌚ PREPROCESSING PIPELINE STEPS COMPLETED:

- ✓ Text cleaning (URLs, special chars, punctuation)
- ✓ Tokenization
- ✓ Stopword removal (English + custom)
- ✓ Lemmatization
- ✓ Empty text removal
- ✓ Feature engineering
- ✓ Data validation
- ✓ Category-wise file generation

=====

= PREPROCESSING COMPLETED SUCCESSFULLY! ✓

=====

=

```
In [17]: print("=*60)
print("FINAL DATA QUALITY CHECK")
print("=*60)

print("\n\n Checking for null values in key columns:")
key_columns = ['cleaned_text', 'health_category', 'combined_text']
for col in key_columns:
    null_count = df_clean[col].isnull().sum()
    print(f" {col}: {null_count} nulls")

print("\n\n Checking for empty strings:")
empty_cleaned = len(df_clean[df_clean['cleaned_text'] == ''])
empty_combined = len(df_clean[df_clean['combined_text'] == ''])
print(f" cleaned_text: {empty_cleaned} empty strings")
print(f" combined_text: {empty_combined} empty strings")

print("\n\n Checking word count distribution:")
very_short = len(df_clean[df_clean['cleaned_word_count'] < 5])
short = len(df_clean[(df_clean['cleaned_word_count'] >= 5) & (df_clean['cleaned_word_count'] < 20)])
medium = len(df_clean[(df_clean['cleaned_word_count'] >= 20) & (df_clean['cleaned_word_count'] < 100)])
long = len(df_clean[df_clean['cleaned_word_count'] >= 100])

print(f" Very short (<5 words): {very_short:,} ({very_short/len(df_clean)*100:.1f}%)")
print(f" Short (5-19 words): {short:,} ({short/len(df_clean)*100:.1f}%)")
print(f" Medium (20-99 words): {medium:,} ({medium/len(df_clean)*100:.1f}%)")
print(f" Long (100+ words): {long:,} ({long/len(df_clean)*100:.1f}%)")

print("\n\n Checking category balance:")
min_category_size = df_clean['health_category'].value_counts().min()
```

```

max_category_size = df_clean['health_category'].value_counts().max()
balance_ratio = min_category_size / max_category_size
print(f" Smallest category: {min_category_size:,} samples")
print(f" Largest category: {max_category_size:,} samples")
print(f" Balance ratio: {balance_ratio:.3f}")
if balance_ratio < 0.3:
    print(f" △ Warning: Significant class imbalance detected!")
else:
    print(f" ✓ Categories are reasonably balanced")

print("\n" + "="*60)
print("DATA QUALITY CHECK COMPLETED! ✓")
print("="*60)

```

=====

FINAL DATA QUALITY CHECK

=====

- ✓ Checking for null values in key columns:
 cleaned_text: 0 nulls
 health_category: 0 nulls
 combined_text: 0 nulls
- ✓ Checking for empty strings:
 cleaned_text: 0 empty strings
 combined_text: 0 empty strings
- ✓ Checking word count distribution:
 Very short (<5 words): 37 (0.2%)
 Short (5-19 words): 2,595 (13.6%)
 Medium (20-99 words): 12,378 (64.9%)
 Long (100+ words): 4,050 (21.2%)
- ✓ Checking category balance:
 Smallest category: 2,519 samples
 Largest category: 2,892 samples
 Balance ratio: 0.871
 ✓ Categories are reasonably balanced

=====

DATA QUALITY CHECK COMPLETED! ✓

=====

In []: