

```

# Cell 1: Import Required Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from collections import Counter
import warnings

warnings.filterwarnings('ignore')
plt.style.use('seaborn-v0_8-darkgrid')
sns.set_palette("husl")

%matplotlib inline

print("All libraries imported successfully!")
All libraries imported successfully!

# Cell 2: Load Kaggle Dataset
data_path =
'../data/raw/Lifestyle_and_Health_Risk_Prediction_Synthetic_Dataset
Kaggle.csv'

df_kaggle = pd.read_csv(data_path)

print("*70)
print(" " * 20 + "KAGGLE DATASET LOADED SUCCESSFULLY")
print("*70)
print(f"Total Records: {len(df_kaggle)}")
print(f"Total Columns: {len(df_kaggle.columns)}")
print(f"\nColumn Names:")
print(df_kaggle.columns.tolist())

=====
KAGGLE DATASET LOADED SUCCESSFULLY
=====
Total Records: 5,000
Total Columns: 12

Column Names:
['age', 'weight', 'height', 'exercise', 'sleep', 'sugar_intake',
'smoking', 'alcohol', 'married', 'profession', 'bmi', 'health_risk']

# Cell 3: Dataset Overview
print("*70)
print(" " * 25 + "DATASET OVERVIEW")
print("*70)

print("\nFirst 10 Rows:")
print(df_kaggle.head(10))

```

```

print("\n" + "*70)
print("Last 5 Rows:")
print(df_kaggle.tail())

print("\n" + "*70)
print("Random Sample (5 rows):")
print(df_kaggle.sample(5))

```

=====

DATASET OVERVIEW

=====

First 10 Rows:

	age	weight	height	exercise	sleep	sugar_intake	smoking	alcohol
married \	56	67	195	low	6.1	medium	yes	yes
yes	69	76	170	high	6.9	high	no	no
no	46	106	153	high	6.6	low	yes	no
no	32	54	186	medium	8.5	medium	no	no
yes	60	98	195	high	8.0	low	no	no
yes	25	96	160	medium	3.8	medium	no	no
yes	78	64	168	medium	9.9	high	no	no
yes	38	76	194	low	6.6	medium	yes	no
yes	56	58	158	medium	9.6	medium	no	no
no	75	94	147	none	8.1	medium	yes	no

	profession	bmi	health_risk
0	office_worker	17.6	high
1	teacher	26.3	high
2	artist	45.3	high
3	artist	15.6	low
4	teacher	25.8	high
5	farmer	37.5	high
6	artist	22.7	high
7	driver	20.2	high
8	office_worker	23.2	low
9	farmer	43.5	high

=====

Last 5 Rows:

```

      age  weight  height exercise   sleep sugar_intake smoking alcohol
\ 4995    42       75     187   medium     7.2        low      no      no
4996    39       86     183     low      7.4        low      no      no
4997    48       56     194   none      7.2        high     yes     yes
4998    34      109     148   medium     9.9        medium    no      no
4999    72       47     149     high      6.3        medium    no      no

      married profession   bmi health_risk
4995      yes      driver  21.4      low
4996      yes      engineer 25.7      high
4997      yes      engineer 14.9      high
4998      yes      doctor   49.8      low
4999      yes      driver   21.2      low
=====

Random Sample (5 rows):
      age  weight  height exercise   sleep sugar_intake smoking alcohol
\ 2786    77       49     186     low      9.8        medium    no      no
4880    20      103     163     low      6.8        medium    no      no
1160    78       76     162   medium     7.1        high     no      no
2135    37      105     165   medium     9.0        low      no      no
1375    72      91      161   medium     6.1        medium    no      no

      married profession   bmi health_risk
2786      no      doctor  14.2      high
4880      no      engineer 38.8      high
1160      yes     artist  29.0      high
2135      yes     doctor  38.6      low
1375      no      farmer  35.1      high

# Cell 4: Data Types and Structure
print("=*70)
print(" " * 25 + "DATA TYPES & INFO")
print("=*70)

print("\nDataFrame Info:")
print(df_kaggle.info())

print("\n" + "*70)

```

```

print("Data Types Summary:")
print(df_kaggle.dtypes)

print("\n" + "*70)
print("Dataset Shape:")
print(f" Rows: {df_kaggle.shape[0]},")
print(f" Columns: {df_kaggle.shape[1]}")

print("\n" + "*70)
print("Memory Usage:")
print(f" {df_kaggle.memory_usage(deep=True).sum() / 1024**2:.2f} MB")

=====
DATA TYPES & INFO
=====

DataFrame Info:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5000 entries, 0 to 4999
Data columns (total 12 columns):
 #   Column            Non-Null Count  Dtype  
 --- 
 0   age               5000 non-null    int64  
 1   weight             5000 non-null    int64  
 2   height             5000 non-null    int64  
 3   exercise           5000 non-null    object  
 4   sleep              5000 non-null    float64 
 5   sugar_intake       5000 non-null    object  
 6   smoking             5000 non-null    object  
 7   alcohol             5000 non-null    object  
 8   married             5000 non-null    object  
 9   profession          5000 non-null    object  
 10  bmi                5000 non-null    float64 
 11  health_risk         5000 non-null    object  
dtypes: float64(2), int64(3), object(7)
memory usage: 468.9+ KB
None

=====
Data Types Summary:
age                  int64
weight               int64
height               int64
exercise             object
sleep                float64
sugar_intake         object
smoking              object
alcohol              object
married              object
profession           object

```

```

bmi           float64
health_risk    object
dtype: object
=====
Dataset Shape:
  Rows: 5,000
  Columns: 12
=====
Memory Usage:
  2.22 MB
# Cell 5: Missing Values Analysis
print("*" * 70)
print(" " * 25 + "MISSING VALUES ANALYSIS")
print("*" * 70)

missing_values = df_kaggle.isnull().sum()
missing_percentage = (missing_values / len(df_kaggle)) * 100

missing_df = pd.DataFrame({
    'Column': missing_values.index,
    'Missing_Count': missing_values.values,
    'Percentage': missing_percentage.values
})

missing_df = missing_df[missing_df['Missing_Count'] > 0].sort_values('Missing_Count', ascending=False)

if len(missing_df) > 0:
    print("\nMissing Values Found:")
    print(missing_df.to_string(index=False))

    # Visualization
    plt.figure(figsize=(12, 6))
    plt.barh(missing_df['Column'], missing_df['Percentage'],
color='coral')
    plt.xlabel('Missing Percentage (%)', fontsize=12)
    plt.ylabel('Column Name', fontsize=12)
    plt.title('Missing Values by Column', fontsize=14,
fontweight='bold')
    plt.gca().invert_yaxis()
    plt.tight_layout()
    plt.savefig('../reports/kaggle_missing_values.png', dpi=300,
bbox_inches='tight')
    plt.show()
else:
    print("\n[] No missing values found!")
    print("    Dataset is complete and ready for analysis.")

```

MISSING VALUES ANALYSIS

□ No missing values found!
Dataset is complete and ready for analysis.

```
# Cell 6: Statistical Summary
print("*" * 70)
print(" " * 25 + "STATISTICAL SUMMARY")
print("*" * 70)

print("\nNumerical Features Summary:")
print(df_kaggle.describe())

print("\n" + "*" * 70)
print("Categorical Features Summary:")
categorical_cols = df_kaggle.select_dtypes(include=['object']).columns
if len(categorical_cols) > 0:
    print(df_kaggle[categorical_cols].describe())
else:
    print("No categorical columns found.")
```

STATISTICAL SUMMARY

Numerical Features Summary:

	age	weight	height	sleep	bmi
count	5000.000000	5000.000000	5000.000000	5000.000000	5000.000000
mean	48.805600	77.359200	171.933200	6.997060	26.84388
std	17.906991	18.714567	15.822225	1.432849	8.25249
min	18.000000	45.000000	145.000000	3.000000	11.40000
25%	34.000000	61.000000	158.000000	6.000000	20.30000
50%	49.000000	77.000000	171.500000	7.000000	26.00000
75%	64.000000	94.000000	186.000000	8.000000	32.40000
max	79.000000	109.000000	199.000000	10.000000	51.40000

Categorical Features Summary:

	exercise	sugar_intake	smoking	alcohol	married	profession
health_risk	5000	5000	5000	5000	5000	5000
count	5000	5000	5000	5000	5000	5000
5000						
unique	4	3	2	2	2	8
2						
top	medium	medium	no	no	yes	student
high						
freq	1985	2511	4023	3741	3018	644
3490						

```

# Cell 7: Target Variable Distribution
print("*" * 70)
print(" " * 20 + "TARGET VARIABLE: health_risk")
print("*" * 70)

# Identify target column (adjust name if different)
target_col = 'health_risk' # Change if your target column has a
                           # different name

if target_col in df_kaggle.columns:

    # Value counts
    risk_counts = df_kaggle[target_col].value_counts()
    risk_percentage = (risk_counts / len(df_kaggle)) * 100

    risk_df = pd.DataFrame({
        'Risk_Level': risk_counts.index,
        'Count': risk_counts.values,
        'Percentage': risk_percentage.values
    })

    print("\nRisk Level Distribution:")
    print(risk_df.to_string(index=False))

    # Visualizations
    fig, axes = plt.subplots(1, 2, figsize=(16, 6))

    # Bar chart
    axes[0].bar(risk_counts.index, risk_counts.values,
                color=sns.color_palette("Set2", len(risk_counts)),
                edgecolor='black')
    axes[0].set_xlabel('Risk Level', fontsize=12)
    axes[0].set_ylabel('Number of Records', fontsize=12)
    axes[0].set_title('Distribution of Health Risk Levels',
                      fontsize=14, fontweight='bold')
    for i, v in enumerate(risk_counts.values):
        axes[0].text(i, v + 10, str(v), ha='center', va='bottom',
                     fontweight='bold')

    # Pie chart
    colors = sns.color_palette("Set2", len(risk_counts))
    wedges, texts, autotexts = axes[1].pie(risk_counts.values,
                                             labels=risk_counts.index,
                                             autopct='%1.1f%%',
                                             startangle=90, colors=colors)
    axes[1].set_title('Percentage Distribution of Risk Levels',
                      fontsize=14, fontweight='bold')

    for autotext in autotexts:
        autotext.set_color('white')

```

```

autotext.set_fontweight('bold')
autotext.set_fontsize(12)

plt.tight_layout()
plt.savefig('../reports/kaggle_risk_distribution.png', dpi=300,
bbox_inches='tight')
plt.show()

# Check class balance
min_class = risk_counts.min()
max_class = risk_counts.max()
balance_ratio = min_class / max_class

print(f"\n\square Class Balance Analysis:")
print(f"    Smallest class: {min_class:,} samples")
print(f"    Largest class: {max_class:,} samples")
print(f"    Balance ratio: {balance_ratio:.3f}")


if balance_ratio < 0.5:
    print("    \u2192 Warning: Class imbalance detected!")
    print("    \u2192 Consider using stratified sampling or SMOTE")
else:
    print("    \u2296 Classes are reasonably balanced")

else:
    print(f"\u2192 Column '{target_col}' not found!")
    print(f"    Available columns: {df_kaggle.columns.tolist()}")
=====

TARGET VARIABLE: health_risk
=====

Risk Level Distribution:
Risk_Level  Count  Percentage
      high    3490      69.8
       low    1510      30.2

```



□ Class Balance Analysis:

Smallest class: 1,510 samples

Largest class: 3,490 samples

Balance ratio: 0.433

⚠ Warning: Class imbalance detected!

→ Consider using stratified sampling or SMOTE

```
# Cell 8: Numerical Features Analysis
print("*" * 70)
print(" " * 25 + "NUMERICAL FEATURES ANALYSIS")
print("*" * 70)

numerical_cols =
df_kaggle.select_dtypes(include=[np.number]).columns.tolist()

# Remove target if it's numerical
if target_col in numerical_cols and df_kaggle[target_col].nunique() <= 10:
    numerical_cols.remove(target_col)

print(f"\nNumerical Features ({len(numerical_cols)}):")
for col in numerical_cols:
    print(f" - {col}")

# Distribution plots
if len(numerical_cols) > 0:
    n_cols = 3
    n_rows = (len(numerical_cols) + n_cols - 1) // n_cols

    fig, axes = plt.subplots(n_rows, n_cols, figsize=(16, 5 * n_rows))
    axes = axes.flatten() if n_rows > 1 else [axes] if n_cols == 1
else axes
```

```

    for idx, col in enumerate(numerical_cols):
        axes[idx].hist(df_kaggle[col].dropna(), bins=30,
color='steelblue',
                           edgecolor='black', alpha=0.7)
        axes[idx].set_xlabel(col, fontsize=11)
        axes[idx].set_ylabel('Frequency', fontsize=11)
        axes[idx].set_title(f'Distribution of {col}', fontsize=12,
fontweight='bold')
        axes[idx].grid(True, alpha=0.3)

        # Add mean line
        mean_val = df_kaggle[col].mean()
        axes[idx].axvline(mean_val, color='red', linestyle='--',
linewidth=2,
                           label=f'Mean: {mean_val:.2f}')
        axes[idx].legend()

# Hide unused subplots
for idx in range(len(numerical_cols), len(axes)):
    axes[idx].axis('off')

plt.tight_layout()
plt.savefig('../reports/kaggle_numerical_distributions.png',
dpi=300, bbox_inches='tight')
plt.show()


print(f"\n\square Saved visualization:
kaggle_numerical_distributions.png")
else:
    print("\nNo numerical features to visualize.")

```

===== NUMERICAL FEATURES ANALYSIS =====

Numerical Features (5):

- age
- weight
- height
- sleep
- bmi



□ Saved visualization: kaggle_numerical_distributions.png

```
# Cell 9: Box Plots for Numerical Features
print("=*70)
print(" " * 20 + "BOX PLOTS - OUTLIER DETECTION")
print("=*70)

if len(numerical_cols) > 0:
    n_cols = 3
    n_rows = (len(numerical_cols) + n_cols - 1) // n_cols

    fig, axes = plt.subplots(n_rows, n_cols, figsize=(16, 5 * n_rows))
    axes = axes.flatten() if n_rows > 1 else [axes] if n_cols == 1
else axes

    for idx, col in enumerate(numerical_cols):
        axes[idx].boxplot(df_kaggle[col].dropna(), vert=True,
patch_artist=True,
color='black'),
                    boxprops=dict(facecolor='lightblue',
medianprops=dict(color='red', linewidth=2),
whiskerprops=dict(color='black'),
capprops=dict(color='black'))
        axes[idx].set_ylabel(col, fontsize=11)
        axes[idx].set_title(f'Box Plot: {col}', fontsize=12,
```

```

        fontweight='bold')
    axes[idx].grid(True, alpha=0.3, axis='y')

    # Hide unused subplots
    for idx in range(len(numerical_cols), len(axes)):
        axes[idx].axis('off')

    plt.tight_layout()
    plt.savefig('../reports/kaggle_boxplots.png', dpi=300,
bbox_inches='tight')
    plt.show()


    print(f"\n\square Saved visualization: kaggle_boxplots.png")

```

=====

BOX PLOTS - OUTLIER DETECTION

=====



\square Saved visualization: kaggle_boxplots.png

```

# Cell 10: Categorical Features Analysis
print("*70)
print(" " * 25 + "CATEGORICAL FEATURES ANALYSIS")
print("*70)

```

```

categorical_cols =
df_kaggle.select_dtypes(include=['object']).columns.tolist()

# Remove target if it's categorical
if target_col in categorical_cols:
    categorical_cols.remove(target_col)

print(f"\nCategorical Features ({len(categorical_cols)}):")
for col in categorical_cols:
    print(f" - {col}")

# Value counts and visualizations
if len(categorical_cols) > 0:
    for col in categorical_cols:
        print(f"\n{'-'*70}")
        print(f"Feature: {col.upper()}")
        print(f"{'-'*70}")

        value_counts = df_kaggle[col].value_counts()
        value_percentage = (value_counts / len(df_kaggle)) * 100

        result_df = pd.DataFrame({
            'Value': value_counts.index,
            'Count': value_counts.values,
            'Percentage': value_percentage.values
        })

        print(result_df.to_string(index=False))

# Visualization
fig, axes = plt.subplots(1, 2, figsize=(14, 5))

# Bar chart
axes[0].bar(range(len(value_counts)), value_counts.values,
            color=sns.color_palette("viridis",
            len(value_counts)), edgecolor='black')
    axes[0].set_xticks(range(len(value_counts)))
    axes[0].set_xticklabels(value_counts.index, rotation=45,
ha='right')
    axes[0].set_ylabel('Count', fontsize=12)
    axes[0].set_title(f'Distribution of {col}', fontsize=14,
fontweight='bold')
    for i, v in enumerate(value_counts.values):
        axes[0].text(i, v + max(value_counts.values)*0.02, str(v),
                    ha='center', va='bottom', fontweight='bold')

# Pie chart
colors = sns.color_palette("viridis", len(value_counts))
wedges, texts, autotexts = axes[1].pie(value_counts.values,

```

```

    labels=value_counts.index,
                                         autopct='%.1f%%',
startangle=90, colors=colors)
        axes[1].set_title(f'Percentage Distribution of {col}',
fontsize=14, fontweight='bold')

        for autotext in autotexts:
            autotext.set_color('white')
            autotext.set_fontweight('bold')

        plt.tight_layout()
        plt.savefig(f'../reports/kaggle_categorical_{col}.png',
dpi=300, bbox_inches='tight')
        plt.show()

        print(f"\nSaved visualization:
kaggle_categorical_{col}.png")
else:
    print("\nNo categorical features to analyze.")

```


CATEGORICAL FEATURES ANALYSIS

Categorical Features (6):

- exercise
- sugar_intake
- smoking
- alcohol
- married
- profession

Feature: EXERCISE


Value	Count	Percentage
medium	1985	39.70
low	1272	25.44
high	1000	20.00
none	743	14.86



□ Saved visualization: kaggle_categorical_exercise.png

Feature: SUGAR_INTAKE

Value	Count	Percentage
medium	2511	50.22
low	1484	29.68
high	1005	20.10





□ Saved visualization: kaggle_categorical_sugar_intake.png

Feature: SMOKING

Value Count Percentage



no	4023	80.46
yes	977	19.54



□ Saved visualization: kaggle_categorical_smoking.png

Feature: ALCOHOL


Value	Count	Percentage
no	3741	74.82
yes	1259	25.18




□ Saved visualization: kaggle_categorical_alcohol.png

Feature: MARRIED

Value	Count	Percentage
yes	3018	60.36
no	1982	39.64




Percentage Distribution of married




□ Saved visualization: kaggle_categorical_married.png

Feature: PROFESSION

Value	Count	Percentage
student	644	12.88
farmer	639	12.78
driver	631	12.62
doctor	630	12.60
artist	626	12.52
engineer	616	12.32
teacher	612	12.24
office_worker	602	12.04



Percentage Distribution of profession



```

□ Saved visualization: kaggle_categorical_profession.png

# Cell 11: Correlation Matrix
print("*"*70)
print(" " * 25 + "CORRELATION ANALYSIS")
print("*"*70)

if len(numerical_cols) > 1:
    # Calculate correlation matrix
    correlation_matrix = df_kaggle[numerical_cols].corr()

    print("\nCorrelation Matrix:")
    print(correlation_matrix)

    # Heatmap
    plt.figure(figsize=(12, 10))
    sns.heatmap(correlation_matrix, annot=True, fmt='.2f',
cmap='coolwarm',
               center=0, square=True, linewidths=1,
cbar_kws={"shrink": 0.8})
    plt.title('Correlation Matrix - Numerical Features', fontsize=16,
fontweight='bold', pad=20)
    plt.tight_layout()
    plt.savefig('../reports/kaggle_correlation_matrix.png', dpi=300,
bbox_inches='tight')
    plt.show()

print(f"\n□ Saved visualization: kaggle_correlation_matrix.png")

# Find highly correlated pairs
print("\n" + "*"*70)
print("Highly Correlated Feature Pairs (|correlation| > 0.7):")
print("*"*70)

high_corr = []
for i in range(len(correlation_matrix.columns)):
    for j in range(i+1, len(correlation_matrix.columns)):
        if abs(correlation_matrix.iloc[i, j]) > 0.7:
            high_corr.append({
                'Feature_1': correlation_matrix.columns[i],
                'Feature_2': correlation_matrix.columns[j],
                'Correlation': correlation_matrix.iloc[i, j]
            })

    if len(high_corr) > 0:
        high_corr_df =
pd.DataFrame(high_corr).sort_values('Correlation',
key=abs,
ascending=False)

```

```
    print(high_corr_df.to_string(index=False))
    print("\n⚠️ High correlation detected - consider feature
selection/engineering")
else:
    print("⚠️ No highly correlated features found")
else:
    print("\n⚠️ Insufficient numerical features for correlation
analysis.")
```

CORRELATION ANALYSIS

Correlation Matrix:

	age	weight	height	sleep	bmi
age	1.000000	0.010220	0.020501	-0.022504	-0.011407
weight	0.010220	1.000000	0.001637	-0.012756	0.784869
height	0.020501	0.001637	1.000000	-0.016776	-0.596046
sleep	-0.022504	-0.012756	-0.016776	1.000000	-0.001544
bmi	-0.011407	0.784869	-0.596046	-0.001544	1.000000

Correlation Matrix - Numerical Features



□ Saved visualization: kaggle_correlation_matrix.png

=====

Highly Correlated Feature Pairs ($|correlation| > 0.7$):

=====

Feature_1	Feature_2	Correlation
weight	bmi	0.784869

△ High correlation detected - consider feature selection/engineering

```
# Cell 12: Features vs Target Analysis
print("=*70)
print(" " * 20 + "FEATURES vs TARGET ANALYSIS")
print("=*70)
```

```

if target_col in df_kaggle.columns:

    # Numerical features vs Target
    if len(numerical_cols) > 0:
        print("\n" + "="*70)
        print("NUMERICAL FEATURES vs RISK LEVEL")
        print("="*70)

        n_cols = 3
        n_rows = (len(numerical_cols) + n_cols - 1) // n_cols

        fig, axes = plt.subplots(n_rows, n_cols, figsize=(16, 5 * n_rows))
        axes = axes.flatten() if n_rows > 1 else [axes] if n_cols == 1 else axes

        for idx, col in enumerate(numerical_cols):
            df_kaggle.boxplot(column=col, by=target_col, ax=axes[idx], patch_artist=True)
            axes[idx].set_xlabel('Risk Level', fontsize=11)
            axes[idx].set_ylabel(col, fontsize=11)
            axes[idx].set_title(f'{col} by Risk Level', fontsize=12, fontweight='bold')
            axes[idx].get_figure().suptitle('')

        # Hide unused subplots
        for idx in range(len(numerical_cols), len(axes)):
            axes[idx].axis('off')

        plt.tight_layout()
        plt.savefig('../reports/kaggle_features_vs_target.png', dpi=300, bbox_inches='tight')
        plt.show()

        print(f"\n\square Saved visualization:\nkaggle_features_vs_target.png")

    # Categorical features vs Target
    if len(categorical_cols) > 0:
        print("\n" + "="*70)
        print("CATEGORICAL FEATURES vs RISK LEVEL")
        print("="*70)

        for col in categorical_cols:
            print(f"\n{col.upper()} vs RISK LEVEL:")
            crosstab = pd.crosstab(df_kaggle[col], df_kaggle[target_col], margins=True)
            print(crosstab)

```


```
# Stacked bar chart
crosstab_no_margins = pd.crosstab(df_kaggle[col],
df_kaggle[target_col], normalize='index') * 100

fig, ax = plt.subplots(figsize=(10, 6))
crosstab_no_margins.plot(kind='bar', stacked=True, ax=ax,
color=sns.color_palette("Set2",
len(df_kaggle[target_col].unique())))
ax.set_xlabel(col, fontsize=12)
ax.set_ylabel('Percentage (%)', fontsize=12)
ax.set_title(f'{col} vs Risk Level (Stacked %)', fontsize=14, fontweight='bold')
ax.legend(title='Risk Level', bbox_to_anchor=(1.05, 1),
loc='upper left')
plt.xticks(rotation=45, ha='right')
plt.tight_layout()
plt.savefig(f'../reports/kaggle_{col}_vs_target.png',
dpi=300, bbox_inches='tight')
plt.show()

print(f"\nSaved visualization:
kaggle_{col}_vs_target.png")
```

FEATURES vs TARGET ANALYSIS

NUMERICAL FEATURES vs RISK LEVEL




□ Saved visualization: kaggle_features_vs_target.png

CATEGORICAL FEATURES vs RISK LEVEL

EXERCISE vs RISK LEVEL:


health_risk	high	low	All
exercise			
high	609	391	1000
low	1065	207	1272
medium	1202	783	1985
none	614	129	743
All	3490	1510	5000



□ Saved visualization: kaggle_exercise_vs_target.png

SUGAR_INTAKE vs RISK LEVEL:


health_risk	high	low	All
sugar_intake			
high	845	160	1005
low	989	495	1484
medium	1656	855	2511
All	3490	1510	5000



□ Saved visualization: kaggle_sugar_intake_vs_target.png

SMOKING vs RISK LEVEL:


	health_risk	high	low	All
smoking				
no		2578	1445	4023
yes		912	65	977
All		3490	1510	5000



□ Saved visualization: kaggle_smoking_vs_target.png

ALCOHOL vs RISK LEVEL:


health_risk	high	low	All
alcohol			
no	2431	1310	3741
yes	1059	200	1259
All	3490	1510	5000



□ Saved visualization: kaggle_alcohol_vs_target.png

MARRIED vs RISK LEVEL:


	health_risk	high	low	All
married				
no		1405	577	1982
yes		2085	933	3018
All		3490	1510	5000



□ Saved visualization: kaggle_married_vs_target.png

PROFESSION vs RISK LEVEL:

health_risk	high	low	All
profession			
artist	430	196	626
doctor	419	211	630
driver	457	174	631
engineer	437	179	616
farmer	442	197	639
office_worker	421	181	602
student	451	193	644
teacher	433	179	612
All	3490	1510	5000



□ Saved visualization: kaggle_profession_vs_target.png

```
# Cell 13: Data Quality Summary
print("=*70)
print(" " * 25 + "DATA QUALITY SUMMARY")
print("*70)

print("\n□ DUPLICATES:")
print("-" * 70)
duplicates = df_kaggle.duplicated().sum()
print(f" Total duplicate rows: {duplicates}")
if duplicates > 0:
    print(f" Percentage: {duplicates/len(df_kaggle)*100:.2f}%")
    print(" ▲ Consider removing duplicates before modeling")
else:
    print(" □ No duplicates found")

print("\n□ UNIQUE VALUES PER COLUMN:")
print("-" * 70)
for col in df_kaggle.columns:
    unique_count = df_kaggle[col].nunique()
    unique_pct = (unique_count / len(df_kaggle)) * 100
    print(f" {col:25}: {unique_count:6,} unique values
({unique_pct:6.2f}%)")

print("\n□ DATA TYPES SUMMARY:")
```

```

print("-" * 70)
dtype_summary = df_kaggle.dtypes.value_counts()
for dtype, count in dtype_summary.items():
    print(f" {str(dtype):15s}: {count} columns")

=====
DATA QUALITY SUMMARY
=====

□ DUPLICATES:
-----
Total duplicate rows: 0
□ No duplicates found

□ UNIQUE VALUES PER COLUMN:
-----
age : 62 unique values ( 1.24%)
weight : 65 unique values ( 1.30%)
height : 55 unique values ( 1.10%)
exercise : 4 unique values ( 0.08%)
sleep : 71 unique values ( 1.42%)
sugar_intake : 3 unique values ( 0.06%)
smoking : 2 unique values ( 0.04%)
alcohol : 2 unique values ( 0.04%)
married : 2 unique values ( 0.04%)
profession : 8 unique values ( 0.16%)
bmi : 383 unique values ( 7.66%)
health_risk : 2 unique values ( 0.04%)

□ DATA TYPES SUMMARY:
-----
object : 7 columns
int64 : 3 columns
float64 : 2 columns

# Cell 14: Feature Engineering Suggestions
print("=*70)
print(" " * 20 + "FEATURE ENGINEERING SUGGESTIONS")
print("=*70)

suggestions = []

# Check for potential BMI calculation
if 'height' in df_kaggle.columns and 'weight' in df_kaggle.columns and
'bmi' not in df_kaggle.columns:
    suggestions.append("□ Create BMI: weight / (height/100)^2")

# Check for age groups
if 'age' in df_kaggle.columns:
    suggestions.append("□ Create age groups: Young (<30), Adult (30-

```

```

50), Senior (>50)")

# Check for binary features
for col in categorical_cols:
    if df_kaggle[col].nunique() == 2:
        suggestions.append(f"\u25a0 Convert '{col}' to binary (0/1)")

# Check for exercise/activity combinations
if 'exercise' in df_kaggle.columns and 'sleep' in df_kaggle.columns:
    suggestions.append("\u25a0 Create lifestyle score: combine exercise + sleep quality")

# Check for substance use flags
substance_cols = [col for col in df_kaggle.columns if any(x in col.lower() for x in ['smoking', 'alcohol', 'drug'])]
if len(substance_cols) > 1:
    suggestions.append(f"\u25a0 Create substance_use_flag: combine {', '.join(substance_cols)}")

if len(suggestions) > 0:
    print("\n\u25a0 Recommended Feature Engineering:")
    for idx, suggestion in enumerate(suggestions, 1):
        print(f"  {idx}. {suggestion}")
else:
    print("\n\u25a0 Dataset appears well-featured. No immediate engineering needed.")

=====

```

FEATURE ENGINEERING SUGGESTIONS

\u25a0 Recommended Feature Engineering:

1. \u25a0 Create age groups: Young (<30), Adult (30-50), Senior (>50)
2. \u25a0 Convert 'smoking' to binary (0/1)
3. \u25a0 Convert 'alcohol' to binary (0/1)
4. \u25a0 Convert 'married' to binary (0/1)
5. \u25a0 Create lifestyle score: combine exercise + sleep quality
6. \u25a0 Create substance_use_flag: combine smoking, alcohol

```

# Cell 15: Final EDA Summary Report
print("*"*80)
print(" " * 25 + "KAGGLE DATASET - EDA SUMMARY")
print("*"*80)

print("\n\u25a0 DATASET OVERVIEW:")
print("-" * 80)
print(f"  Total Records: {len(df_kaggle):,}")
print(f"  Total Features: {len(df_kaggle.columns)}")
print(f"  Numerical Features: {len(numerical_cols)}")
print(f"  Categorical Features: {len(categorical_cols)}")

```

```

print(f" Memory Usage: {df_kaggle.memory_usage(deep=True).sum() / 1024**2:.2f} MB")

print("\n DATA QUALITY:")
print("-" * 80)
missing_total = df_kaggle.isnull().sum().sum()
print(f" Missing Values: {missing_total} ({missing_total/(len(df_kaggle)*len(df_kaggle.columns))*100:.2f}%)")
print(f" Duplicate Rows: {df_kaggle.duplicated().sum()}")

if target_col in df_kaggle.columns:
    print(f"\n TARGET VARIABLE ({target_col}):")
    print("-" * 80)
    print(f" Unique Values: {df_kaggle[target_col].nunique()}")
    for value, count in df_kaggle[target_col].value_counts().items():
        pct = count / len(df_kaggle) * 100
        print(f" {value:15s}: {count:6,} ({pct:5.1f}%)")

print("\n VISUALIZATIONS GENERATED:")
print("-" * 80)
visualizations = [
    "kaggle_risk_distribution.png",
    "kaggle_numerical_distributions.png",
    "kaggle_boxplots.png",
    "kaggle_correlation_matrix.png",
    "kaggle_features_vs_target.png"
]
for viz in visualizations:
    print(f" {viz}")

print("\n NEXT STEPS:")
print("-" * 80)
print(" 1. Handle missing values (if any)")
print(" 2. Remove duplicates (if any)")
print(" 3. Engineer new features (see suggestions)")
print(" 4. Encode categorical variables")
print(" 5. Split data (train/test)")
print(" 6. Scale numerical features")
print(" 7. Train baseline models")

print("\n" + "*80")
print(" " * 25 + "KAGGLE EDA COMPLETED SUCCESSFULLY! " + "*80")

```

=====

KAGGLE DATASET - EDA SUMMARY

=====

□ DATASET OVERVIEW:

```
-----  
Total Records: 5,000  
Total Features: 12  
Numerical Features: 5  
Categorical Features: 6  
Memory Usage: 2.22 MB
```

□ DATA QUALITY:

```
-----  
Missing Values: 0 (0.00%)  
Duplicate Rows: 0
```

□ TARGET VARIABLE (health_risk):

```
-----  
Unique Values: 2  
high : 3,490 ( 69.8%)  
low : 1,510 ( 30.2%)
```

□ VISUALIZATIONS GENERATED:

```
-----  
□ kaggle_risk_distribution.png  
□ kaggle_numerical_distributions.png  
□ kaggle_boxplots.png  
□ kaggle_correlation_matrix.png  
□ kaggle_features_vs_target.png
```

□ NEXT STEPS:

- ```

1. Handle missing values (if any)
2. Remove duplicates (if any)
3. Engineer new features (see suggestions)
4. Encode categorical variables
5. Split data (train/test)
6. Scale numerical features
7. Train baseline models
```

```
=====
=====
KAGGLE EDA COMPLETED SUCCESSFULLY! □
=====
=====
```